A Truthful Training-Run Monitor for High-Stakes Decisions

Live App: https://programmerbaddie.github.io/soup-training-monitor/ | **Repo:** https://github.com/programmerbaddie/soup-training-monitor

1. Data Truthfulness: Can You Show It to Someone Stopping a $500 Run?

When an engineer decides whether to terminate a $500 GPU training run, a chart that looks clean and a chart that is truthful are opposites. Standard dashboards employ naive striding (sampling every *N*th point) or triangular area decimation (LTTB) that silently drops isolated 1-point gradient blowups. If step 524,288 explodes to loss 942.5 while adjacent points remain at 1.34, naive charts smooth the line, misleading the engineer into wasting hundreds of dollars.

The Solution: We implement **Min-Max Bucket Downsampling**. The 1,000,000-point timeline is partitioned into screen pixel bins (2,000 buckets). Within every bucket, the algorithm strictly extracts both the minimum and maximum points and places them in chronological order of their original steps. This mathematically guarantees that any outlier peak is retained at 100% fidelity at every zoom level without temporal inversion or false zigzags.

2. Automated Test Proofs & Limitations

The chart engine is verified via 5 deterministic automated tests (Vitest v5.1.0):

• **Spike Invariance:** Tested 1M series downsampled across 500, 1000, 2000, and 5000 buckets. Min-Max preserved exact peak 942.5 at step 524,288 in all runs. Under identical resolutions, naive striding failed, reporting max ~1.39.

• **NaN / Infinity Integrity:** Confirmed that NaN values emit explicit null line breaks (spanGaps: false). NaN is never converted to 0.0 or silently interpolated.

• **Monotonic Ordering:** Verified that distributed DDP out-of-order logs and duplicate collisions are sorted and deduplicated, eliminating visual loopback artifacts.

• **Log Scale Guard:** Verified that non-positive values (≤0) trigger an epsilon clamp (1e-6) and surface an explicit alert banner.

Limitations (What automated tests would NOT catch): Automated numerical assertions cannot verify human visual perception thresholds (a 1px vertical line might be missed by an eye on high-DPI screens without point markers), nor can they detect NTP wall-clock drift when step indices increase monotonically but real-time logging lagged.

Automated Test Proof	Condition Tested	Observed Result	Status
1M Spike Retention	Step 524,288 loss=942.5 across 4 zoom resolutions	Max=942.5 preserved in 100% of buckets	**PASS (379ms)**
NaN Clamping Rejection	Loss stream containing deliberate NaNs	Emits null breaks; min never poisoned to 0.0	**PASS (1ms)**
DDP Collision Sanitation	Out-of-order steps and duplicate step retries	Timeline strictly monotonically increasing	**PASS (1ms)**
Log Scale Zero Protection	Loss values <= 0 in logarithmic scale mode	Guards -Inf, flags 2 non-positive points	**PASS (1ms)**

3. Performance in Numbers (1,000,000 Points · CPU Throttling ×4)

To achieve instantaneous interaction under heavy CPU throttling, we selected **uPlot** (Canvas-based) paired with typed Float64Array data structures and Min-Max downsampling.

Performance Metric	Before Optimization (SVG / Raw Canvas)	After Optimization (uPlot + Min-Max)	Gain
1M Points First Render	>18,000 ms (DOM choke / Tab Crash)	**413 ms** (Generation 215ms + Downsample 171ms + Draw 10ms)	**~43x faster**
Zoom / Pan Frame Response	120–180 ms (Un-downsampled Canvas lag)	**4.8 ms** (Under 4x CPU throttle, rock-solid 60 fps)	**35x faster**
Hover Tooltip Tracking	45 ms (Frequent frame drops)	**1.2 ms** (Zero perceptible latency)	**37x faster**
Tab Heap Memory (1M pts)	>850 MB (High garbage collection pauses)	**34.5 MB** (30.5 MB raw payload + 4 MB render buffer)	**96% reduction**

4. Silent Failures Audit: What Looks Fine But Is Wrong

We audited the system against five critical failure modes where dashboards present deceptively calm interfaces:
1. Request Race on Fast Run Switching: Rapidly clicking Run 1 → Run 20 (1M points) → Run 2 causes Run 20's delayed computation to finish last, overwriting Run 2 with the wrong telemetry. *Audit & Fix:* Implemented monotonic request tokens (activeRequestIdRef). Any resolution whose token does not match the latest selection is discarded before state mutation.
2. Log Scale with Zeros and Negatives: Computing log10(0) returns -Infinity, while negative contrastive margins produce NaN. Standard charts drop these points silently. *Audit & Fix:* Filter sweeps detect non-positive values, clamp them to an epsilon floor (1e-6) for continuous visual rendering, and display a prominent warning banner stating the exact count.
3. Smoothing Hiding Catastrophic Divergence: Running Exponential Moving Averages (EMA) over noisy gradients dampens sudden upward inflection points. *Audit & Fix:* Min-Max downsampling replaces EMA for default views, ensuring that raw extreme variance is always visible.
4. Stale Data & Background Desync: Single-page dashboards frequently display cached telemetry without notifying the user that polling halted. *Audit & Fix:* Run cards carry explicit step timestamps and telemetry status indicators, making stale streams immediately distinguishable from live runs.
5. Timezone & Locale Inconsistencies: Parsing raw datetime strings via client-local formatters shifts epoch boundaries across distributed team members. *Audit & Fix:* All timestamps are serialized in ISO-8601 UTC and rendered relative to training step indices rather than client wall clocks.

5. What Surprised Us & What Still Concerns Us

What surprised us: How easily standard industry charting paradigms lie without raising errors. While constructing the Naive Stride baseline, observing a 942.5 loss explosion at step 524,288 disappear into a calm 1.34 line with zero console warnings was striking. A machine learning team relying on default downsampling would have wasted hours of cluster compute without realizing their model was diverging.
What still concerns us: Display aspect ratios versus human eye perception. Even when extrema are mathematically preserved in the dataset, a 1-pixel-wide vertical spike on a 4K display occupies fractions of a millimeter. If an engineer glances at the monitor while scrolling, their eye can easily overlook it. In production, single-step outlier spikes must not rely solely on line geometry—they require automated anomaly threshold detection and persistent UI warning badges.

6. AI Tools Disclosure

• **Where Used:** AI assistance was utilized for boilerplate UI scaffolding (Tailwind CSS utility styling, Lucide icon wiring, and TypeScript type declarations).
• **Suggestions Accepted:** Accepted Vite + React 19 configuration and basic layout container structure.
• **What We Verified, Changed, or Redid:** AI initially proposed Recharts and standard array slicing for downsampling. Both were rejected: Recharts crashed the tab at 1,000,000 points (>850MB memory), and standard slicing skipped the spike. We independently designed and implemented the Min-Max Bucket Downsampling algorithm, configured uPlot's canvas coordinate mappings, wrote the deterministic Vitest proof suite, and implemented the active-token race condition guard.

7. Final Verdict: SHIP or DON'T SHIP?

VERDICT: SHIP (with 1 immediate production enhancement)
The system fulfills the core requirement: **uncompromising data truthfulness** under production-scale data volumes. The monitor handles 1,000,000 points smoothly at 60 fps, single-step spikes are mathematically guaranteed to appear at every zoom level, NaNs break the line explicitly, and race conditions are mitigated.
What to change first: Move the synthetic generation and Min-Max downsampling pass from the main thread into a dedicated **Web Worker** using Transferable ArrayBuffer. Although 413 ms is already fast, worker offloading will guarantee that the UI thread maintains 0 ms input latency even when loading multi-million-point runs.